Zhangfei Yang

M.S. Student in Computer Science
The George Washington University
Washington, DC

Email: yzf5140@gmail.com
Home Page: zhangfeiy.github.io

Education

M.S. The George Washington University 2025 - 2027
Computer Science Washington, DC
GPA: 4.0/4.0 | SEAS Achievement Scholarship

B.S. Zhongnan University of Economics and Law 2019 - 2023
Information Systems Wuhan, China
GPA: 3.28/4.0 | National Scholarship

Research Interests

Trustworthy AI, Large Language Model Alignment, Model Behavior and Evaluation, AI Safety

Research Experience

Research Assistant Aug. 2026 - Current
The George Washington University | *Advisor: Prof. Shi Feng*
Review studies of LLM specification gaming and reward seeking, focusing on monitoring cues that can also act as instructions.

AiChatbox | Evidence Grounding in Medical Question Answering Oct. 2026 - Current
Research collaboration with Yi Zhu | Evaluation questions from Meghna Shukla (Wayne State University College of Nursing)
Developed a retrieval-augmented QA system for sickle cell disease education. Designed a claim-level evaluation to test whether cited passages support treatment claims and whether answers preserve qualifications that change care guidance.

Publications

Training-Free Adaptive 360° Video Streaming via Semantic Potential Fields
Aizierjiang Aiersilan, **Zhangfei Yang**
In Proceedings of the 35th International Conference on Computer Communications and Networks (ICCCN), 2026. arXiv:2603.20999.

Internal Coherence Maximization under Decentralized Data
Zhangfei Yang, Aizierjiang Aiersilan | *Oral presentation (presenter)*
Accepted to the AAAI 2026 Fall Symposium on Trustworthy Agentic Systems (TAS 2026).

The Variance Problem in Human Oversight of Clinical AI: Measuring Four Safeguards Across Nine Reviewers on a Nurse-Facing Deterioration Testbed
Zhangfei Yang, Yuheng Zhao, Zhentao Fan
Accepted to the AAAI 2026 Fall Symposium on Agentic and Trustworthy AI for Health and the Global AI-Ready Nursing Workforce (AT-AI4H-NW 2026).

Moral Foundations Reward Modeling for Pluriversal LLM Alignment
Zhangfei Yang | *Oral presentation (presenter)*
Accepted to the AAAI 2026 Fall Symposium "Aligning with Whom? Pluriversal and Paraconsistent Frameworks for Value Alignment in AI Systems."

Reinforcement Learning from Ecological Feedback (RLEF): A Framework for Aligning Language-Model Planning with Planetary Boundaries
Zhangfei Yang, Aizierjiang Aiersilan | *Oral presentation (presenter)*
Accepted to the AAAI 2026 Fall Symposium on Planning for a Better Planet, 2026.

From Delegates to Trustees: Aligning LLM Planners with Absent Stakeholders in Climate Decision-Making
Zhangfei Yang | *Oral presentation (presenter)*
Accepted to the AAAI 2026 Fall Symposium on Planning for a Better Planet, 2026.

Manuscripts

Persona-Induced Information Asymmetry in Large Language Models
Zhangfei Yang
Manuscript preparing for submission to ACL Rolling Review (ARR).

Metacognitive Alignment
Zhangfei Yang, Lintao Mei
Submitted to the AAAI Fall Symposium Series (AAAI FSS), 2026.

Professional Service

PC Member AAAI Fall Symposium Series 2026
External Reviewer AAAI/ACM Conference on AI, Ethics, and Society (AIES) 2026