Zhangfei Yang

M.S. Student in Computer Science
The George Washington University
Washington, DC

Email: yzf5140@gmail.com
Home Page: zhangfeiy.github.io

Education

M.S. The George Washington University 2025 - 2027
Computer Science Washington, DC
GPA: 4.0/4.0 | SEAS Achievement Scholarship

B.S. Zhongnan University of Economics and Law 2019 - 2023
Information Systems Wuhan, China
GPA: 3.28/4.0 | National Scholarship

Research Interests

Trustworthy AI, Large Language Model Alignment, Model Behavior and Evaluation, AI Safety

Research Experience

Research Assistant Aug. 2026 - Current
The George Washington University | *Advisor: Prof. Shi Feng*
Review studies of LLM specification gaming and reward seeking, focusing on monitoring cues that can also act as instructions.

AiChatbox | Evidence Grounding in Medical Question Answering Oct. 2026 - Current
Research collaboration with Yi Zhu | Evaluation questions from Meghna Shukla (Wayne State University College of Nursing)
Developed a retrieval-augmented QA system for sickle cell disease education. Designed a claim-level evaluation to test whether cited passages support treatment claims and whether answers preserve qualifications that change care guidance.

Publications

Training-Free Adaptive 360° Video Streaming via Semantic Potential Fields
Aizierjiang Aiersilan, **Zhangfei Yang**
In Proceedings of the 35th International Conference on Computer Communications and Networks (ICCCN), 2026. arXiv:2603.20999.

Internal Coherence Maximization Under Decentralized Data
Zhangfei Yang, Aizierjiang Aiersilan | *Oral presentation (presenter)*
Proceedings of the AAAI Symposium Series, 10(1), 1132–1139 (2026) · Trustworthy Agentic Systems

The Variance Problem in Human Oversight of Clinical AI: Measuring Four Safeguards Across Nine Reviewers on a Nurse-Facing Deterioration Testbed
Zhangfei Yang, Yuheng Zhao, Zhentao Fan
Proceedings of the AAAI Symposium Series, 10(1), 1281–1288 (2026) · AT-AI4H-NW

Moral Foundations Reward Modeling for Pluriversal LLM Alignment
Zhangfei Yang | *Oral presentation (presenter)*
Proceedings of the AAAI Symposium Series, 10(1), 48–55 (2026) · Aligning with Whom?

Reinforcement Learning from Ecological Feedback (RLEF): A Framework for Aligning Language-Model Planning with Planetary Boundaries
Zhangfei Yang, Aizierjiang Aiersilan | *Oral presentation (presenter)*
Proceedings of the AAAI Symposium Series, 10(1), 566–574 (2026) · Planning for a Better Planet

From Delegates to Trustees: Aligning LLM Planners with Absent Stakeholders in Climate Decision-Making
Zhangfei Yang | *Oral presentation (presenter)*
Proceedings of the AAAI Symposium Series, 10(1), 557–565 (2026) · Planning for a Better Planet

Manuscripts

Persona-Induced Information Asymmetry in Large Language Models
Zhangfei Yang
Manuscript preparing for submission to ACL Rolling Review (ARR).

Professional Service

PC Member AAAI Fall Symposium Series 2026
External Reviewer AAAI/ACM Conference on AI, Ethics, and Society (AIES) 2026